


ARCHITCTUR2

System architecture diagram



OSI model layer breakdown of your protocols

Network Layer: IP (IPv4/IPv6) handles routing packets, Ensures that all data reaches the correct destination.


Transport Layer: TCP for HTTP/1 and HTTP/2

Session Layer: HTTP/1, HTTP/2 manage sessions and connection lifetimes.

Presentation Layer: Protocol Buffers , JSON formatting.

Application Layer: gRPC service methods, HTTP/1 requests, Redis Pub/Sub, SSE for real-time updates.

Sequence diagrams for key flows



Why you chose specific patterns

Server-to-Server gRPC

- For high-frequency metric ingestion system, gRPC is used for server-to-server communication. It allows remote calls to feel like local function invocations while leveraging HTTP/2's multiplexed

streams, low-latency communication, and efficient binary serialization. This makes gRPC ideal for high-throughput, low-overhead data transfer between servers.

dashboard-to-Application http1

- http/1 is preferred for real-time browser updates due to its simplicity, native support, and lower overhead compared to gRPC-Web which is heavy and more complex .

Decoupled Aggregation Layer

- Metrics are sent to a separate aggregation service, using Redis as a temporary store.
- This decoupling reduces latency for each request, thereby making the system more responsive and improving throughput.
- publish-subscribe pattern avoids constant polling and ensures dashboards receive near real-time updates.

Batching for Efficiency

- Metrics are **batched** before sending, which **reduces the number of messages and network load** while maintaining timely updates.
 - **Problem with sending individually:** Each JSON metric ≈ 100 bytes; at 10,000 metrics/sec for 500 dashboards, this would be 500 MB/s (~4 Gbps), which is highly resource-intensive and risks saturating server and network resources.
-

Design decisions and trade-offs

1. gRPC for Server-to-Server Communication:

Browsers cannot natively use full gRPC streams, so an extra layer is needed for client(engineer) communication.

2. Decoupled Aggregation Layer

adding layer for aggregation requires monitoring Redis performance.

3. Batching Metrics

batching slightly increases the latency per metric.